



Convolutional Neural Networks



Deep Learning

- Quick Review of Previous Topics and NNs
- New Theory Topics
- Discuss Famous MNIST Data Set
- Solve MNIST with a “normal” NN
- Learn about CNN
- Solve MNIST with CNN
- CNN Exercise and Solutions Afterwards



Deep Learning

Let's get started!



Quick Review



Deep Learning

**Let's quickly review what
we've covered so far!**



Deep Learning

- Single Neuron
- We now understand how to perform a calculation in a neuron
 - $w \cdot x + b = z$
 - $a = \sigma(z)$



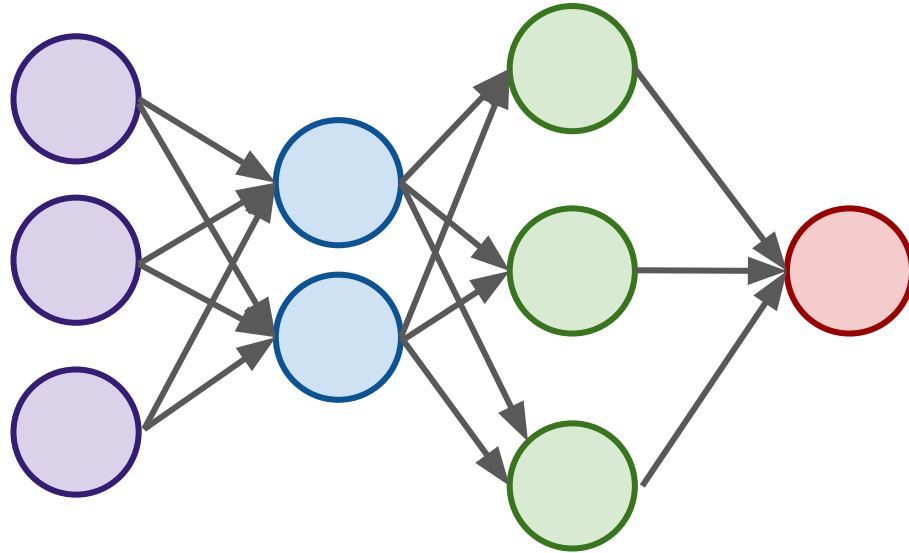
Deep Learning

- Activation Functions
 - Perceptrons
 - Sigmoid
 - Tanh
 - ReLU
 - (We'll discuss other functions later on)



Deep Learning

- Connecting Neurons to create a network





Deep Learning

- Neural Network
 - Input Layer
 - Hidden Layers
 - Output Layer
- More layers → More Abstraction



Deep Learning

- To “learn” we need some measurement of error.
- We use a Cost/Loss Function
 - Quadratic
 - Cross-Entropy



Deep Learning

- Once we have the measurement of error, we need to minimize it by choosing the correct weight and bias values.
- We use gradient descent to find the optimal values.



Deep Learning

- We can then backpropagate the gradient descent through multiple layers, from the output layer back to the input layer.
- We also know of dense layers, and later on we'll introduce softmax layer.



Deep Learning

- We can then backpropagate the gradient descent through multiple layers, from the output layer back to the input layer.
- We also know of dense layers, and later on we'll introduce softmax layer.



Deep Learning

- Very quick overview of what we know so far!
- We still need to learn a bit more theory before diving into Convolutional Neural Networks!



New Theory Topics



Deep Learning

- We've reviewed the basics, but there are still some theory components we haven't covered yet...



Deep Learning

- Initialization of Weights Options
 - Zeros
 - No Randomness
 - Not a great choice
 - Random Distribution Near Zero
 - Not Optimal
 - Activation Functions Distortion



Deep Learning

- Initialization of Weights Options
 - Xavier (Glorot) Initialization
 - Uniform / Normal
 - Draw weights from a distribution with zero mean and a specific variance

$$\text{Var}(W) = \frac{1}{n_{\text{in}}}$$



Deep Learning

- Xavier Initialization

$$Y = W_1 X_1 + W_2 X_2 + \cdots + W_n X_n$$

$$\text{Var}(W_i X_i) = E[X_i]^2 \text{Var}(W_i) + E[W_i]^2 \text{Var}(X_i) + \text{Var}(W_i) \text{Var}(X_i)$$

$$\text{Var}(W_i X_i) = \text{Var}(W_i) \text{Var}(X_i)$$



Deep Learning

- Xavier Initialization

$$\text{Var}(Y) = \text{Var}(W_1X_1 + W_2X_2 + \cdots + W_nX_n) = n\text{Var}(W_i)\text{Var}(X_i)$$

$$\text{Var}(W_i) = \frac{1}{n} = \frac{1}{n_{\text{in}}}$$

$$\text{Var}(W_i) = \frac{2}{n_{\text{in}} + n_{\text{out}}}$$



Deep Learning - Gradient Descent

- Learning Rate - defines the step size during gradient descent
- Batch Size - batches allow us to use stochastic gradient descent.
 - Smaller → less representative of data
 - Larger → longer training time



Deep Learning - Gradient Descent

- Second-Order Behavior of the gradient descent allows us to adjust our learning rate based off the rate of descent
 - AdaGrad
 - RMSProp
 - Adam



Deep Learning - Gradient Descent

- This allows us to start with larger steps and then eventually go to smaller step sizes.
- Adam allows this change to happen automatically.



Deep Learning

- Unstable / Vanishing Gradients
 - As you increase the number of layers in a network, the layers towards the input will be affected less by the error calculation occurring at the output as you go backwards through the network



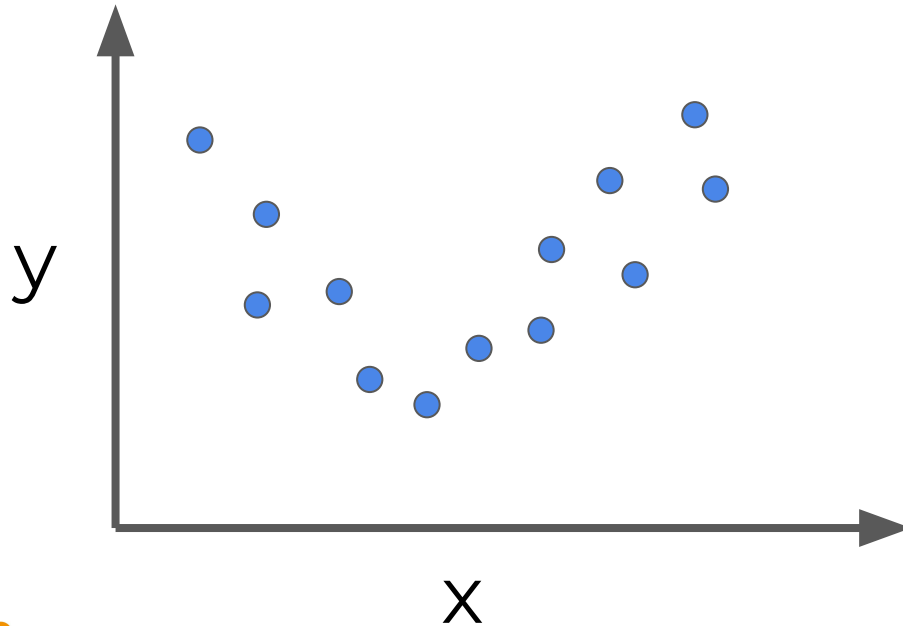
Deep Learning

- Unstable / Vanishing Gradients
 - Initialization and Normalization will help us mitigate these issues.
 - We'll discuss vanishing gradients again in more detail when discussing Recurrent Neural Networks.



Deep Learning

- Overfitting vs Underfitting a Model

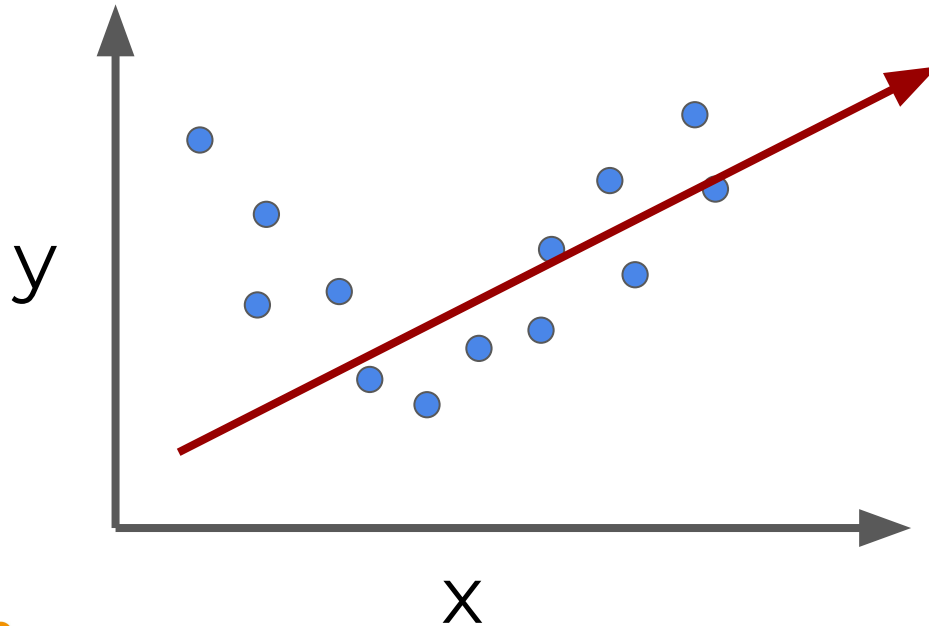


Training
Data



Deep Learning

- Model Underfitting

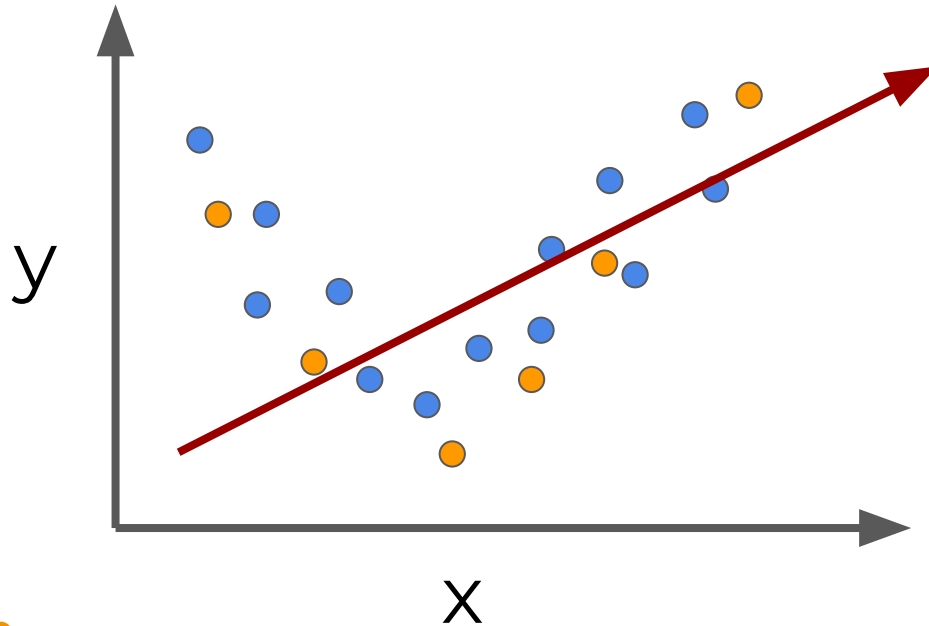


Fitted Model
on Training
Data



Deep Learning

- Model Underfitting




Get error on
test data.



Deep Learning

- Model Underfitting



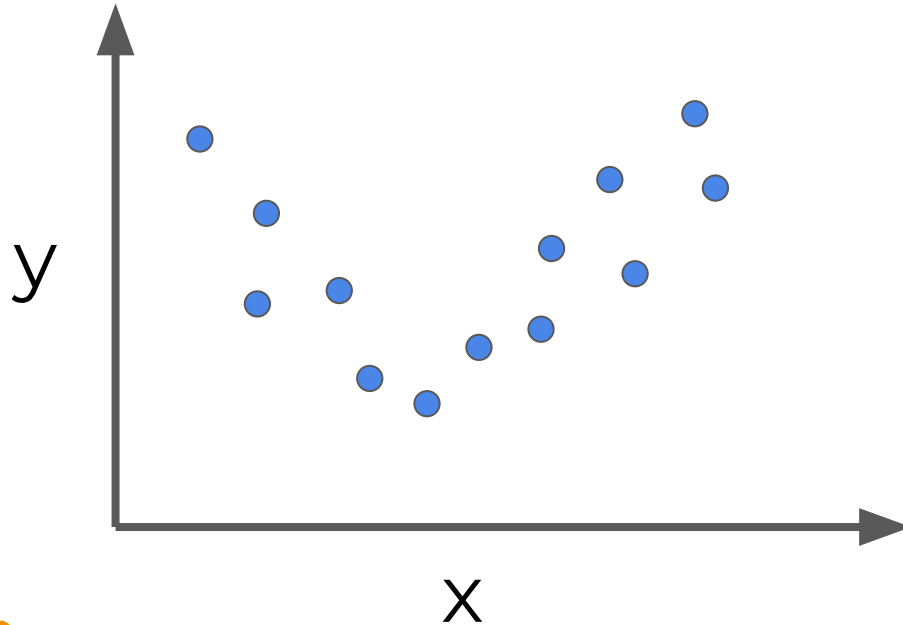
Get error on test data.

High error on training and test data



Deep Learning

- Overfitting a Model

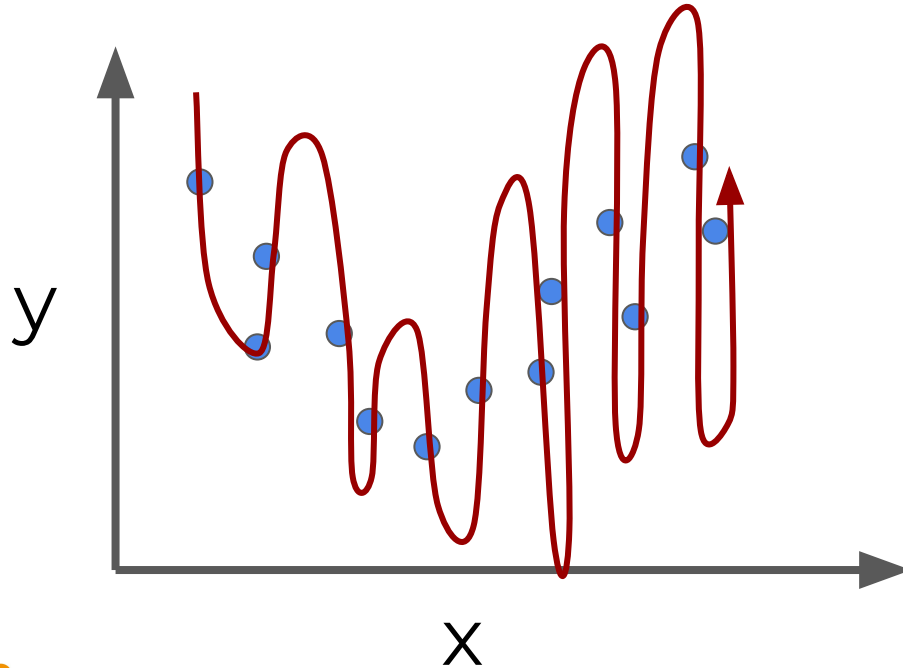


Training
Data



Deep Learning

- Overfitting a Model

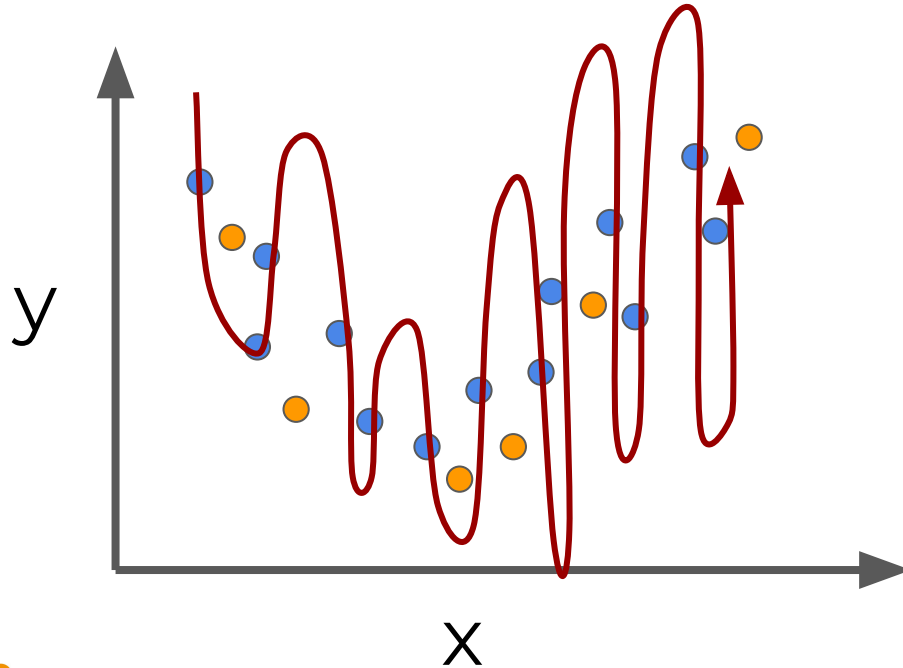


Very low
error on
training
data!



Deep Learning

- Overfitting a Model

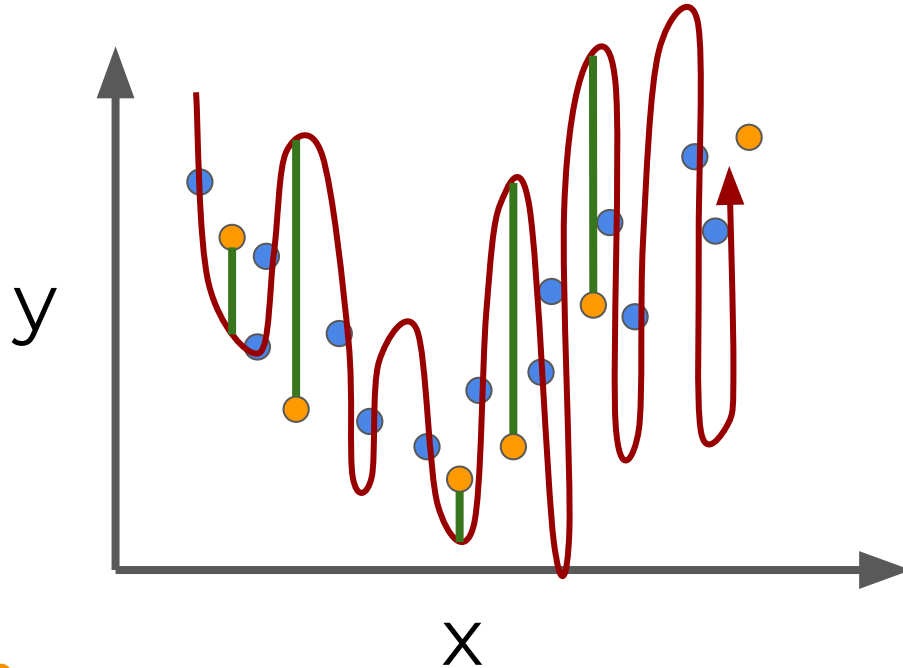


But large
error on test
data.



Deep Learning

- Overfitting a Model




But large
error on test
data.



Deep Learning

- Need to strike a balance





Deep Learning

- With potentially hundreds of parameters in a deep learning neural network, the possibility of overfitting is very high!
- There are a few ways to help mitigate this issue.



Deep Learning

- L1/L2 Regularization
 - Adds a penalty for larger weights in the model
 - Not unique to neural networks



Deep Learning

- Dropout
 - Unique to neural networks
 - Remove neurons during training randomly
 - Network doesn't over rely on any particular neuron



Deep Learning

- Expanding Data
 - Artificially expand data by adding noise, tilting images, adding low white noise to sound data, etc...



Deep Learning

- We still have more theory to learn, such as pooling layers, convolutional layers, etc...
- But we'll wait until we begin to build CNNs to cover those!
- Let's explore the famous MNIST data set, a must know for CNN!



MNIST Data



Deep Learning

- A classic data set in Deep Learning is the MNIST data set.
- Let's quickly cover some basics about it since we'll be using it quite frequently during this section of the course!



Deep Learning

- Fortunately this data is easy to access with TensorFlow. TF has:
 - 55,000 training images
 - 10,000 test images
 - 5,000 validation images



Deep Learning

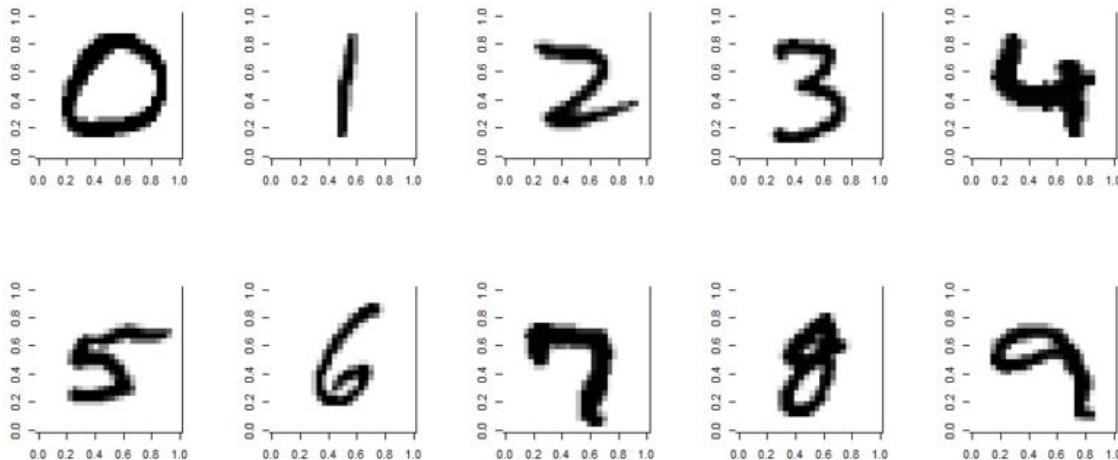
- The MNIST data set contains handwritten single digits from 0 to 9





Deep Learning


- A single digit image can be represented as an array





Deep Learning

- Specifically, 28 by 28 pixels



21


0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.6	.8	0	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0	0
0	0	0	0	0	0	.5	1	.4	0	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0	0
0	0	0	0	0	0	0	1	.7	0	0	0	0	0	0
0	0	0	0	0	0	0	1	1	0	0	0	0	0	0
0	0	0	0	0	0	0	.9	1	.1	0	0	0	0	0
0	0	0	0	0	0	0	.3	1	.1	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0



Deep Learning

- The values represent the grayscale

in




21

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.6	.8	0	0	0	0	0	0
0	0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	0	.5	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	0	1	.7	0	0	0	0	0
0	0	0	0	0	0	0	0	1	1	0	0	0	0	0
0	0	0	0	0	0	0	0	.9	1	.1	0	0	0	0
0	0	0	0	0	0	0	0	.3	1	.1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0



Deep Learning

- We can flatten this array to a 1-D vector of 784 numbers. Either (781,1) or (1,781) is fine, as long as the dimensions are consistent.






Deep Learning

- Flattening out the image ends up removing some of the 2-D information, such as the relationship of a pixel to its neighboring pixels.
- For now, we'll ignore this, but come back to it later when we discuss CNN in depth!



Deep Learning

- We can think of the entire group of the 55,000 images as a tensor (an n-dimensional array)





Deep Learning

- For the labels we'll use One-Hot Encoding.
- This means that instead of having labels such as "One", "Two", etc... we'll have a single array for each image.




Deep Learning

- The label is represented based off the index position in the label array.
- The corresponding label will be a 1 at the index location and zero everywhere else.
- For example, 4 would have this label array:
 - `[0,0,0,0,1,0,0,0,0,0]`



Deep Learning

- As a result, the labels for the training data ends up being a large 2-d array (10,55000):





MNIST Data

“Basic” Approach



Deep Learning

- Before we dive into using CNN on the MNIST data set we'll use a more basic Softmax Regression Approach.
- Let's quickly go over this method (it's very similar to the methods used in the previous sections)



Deep Learning

- A Softmax Regression returns a list of values between 0 and 1 that add up to one.
- We can use this as a list of probabilities!

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$



Deep Learning

- We'll use Softmax as our activation function.

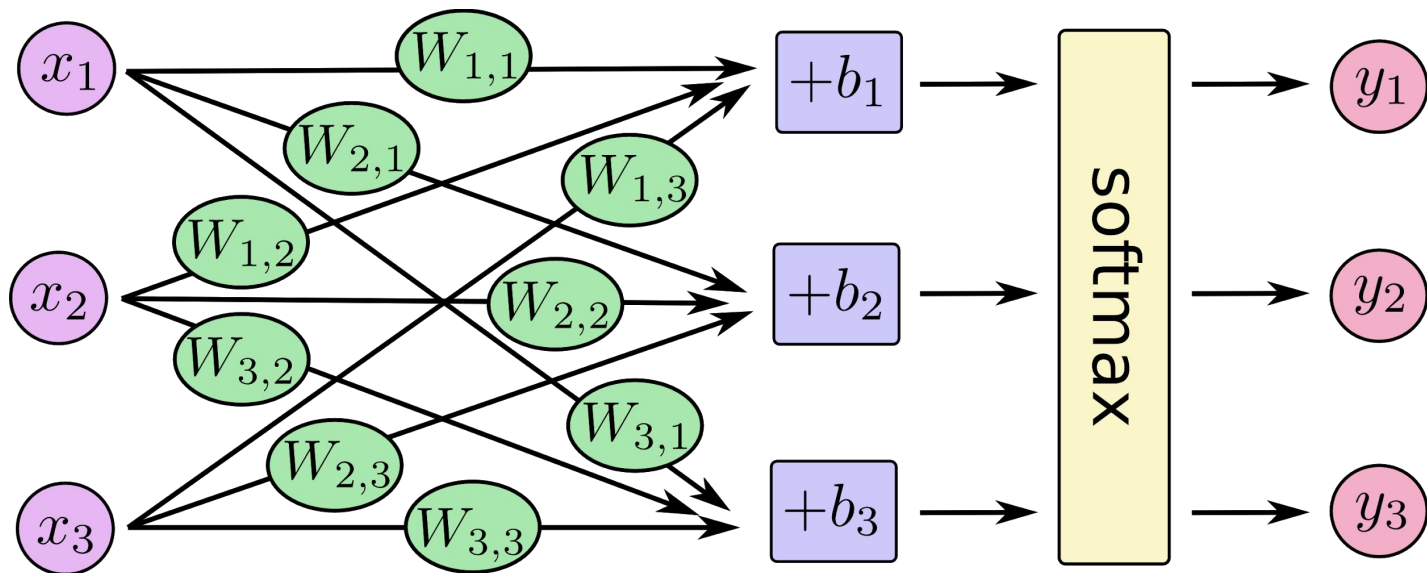
$$z_i = \sum_j W_{i,j} x_j + b_i$$

$$y = \text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$



Deep Learning

Our Network





Deep Learning

As an equation

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \text{softmax} \left(\begin{bmatrix} W_{1,1}x_1 + W_{1,2}x_2 + W_{1,3}x_3 + b_1 \\ W_{2,1}x_1 + W_{2,2}x_2 + W_{2,3}x_3 + b_2 \\ W_{3,1}x_1 + W_{3,2}x_2 + W_{3,3}x_3 + b_3 \end{bmatrix} \right)$$



Deep Learning

As an equation

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \text{softmax} \left(\begin{bmatrix} W_{1,1} & W_{1,2} & W_{1,3} \\ W_{2,1} & W_{2,2} & W_{2,3} \\ W_{3,1} & W_{3,2} & W_{3,3} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} \right)$$



Deep Learning

- Let's implement this with Python and TensorFlow!



Convolutional Neural Networks



Deep Learning

- We just solved the MNIST task with a very simple linear approach.
- Let's explore a much better approach using Convolutional Neural Network.




Deep Learning

- Just like the simple perceptron, CNNs also have their origins in biological research.
- Hubel and Wiesel studied the structure of the visual cortex in mammals, winning a Nobel Prize in 1981.



Deep Learning

- Their research revealed that neurons in the visual cortex had a small local receptive field.





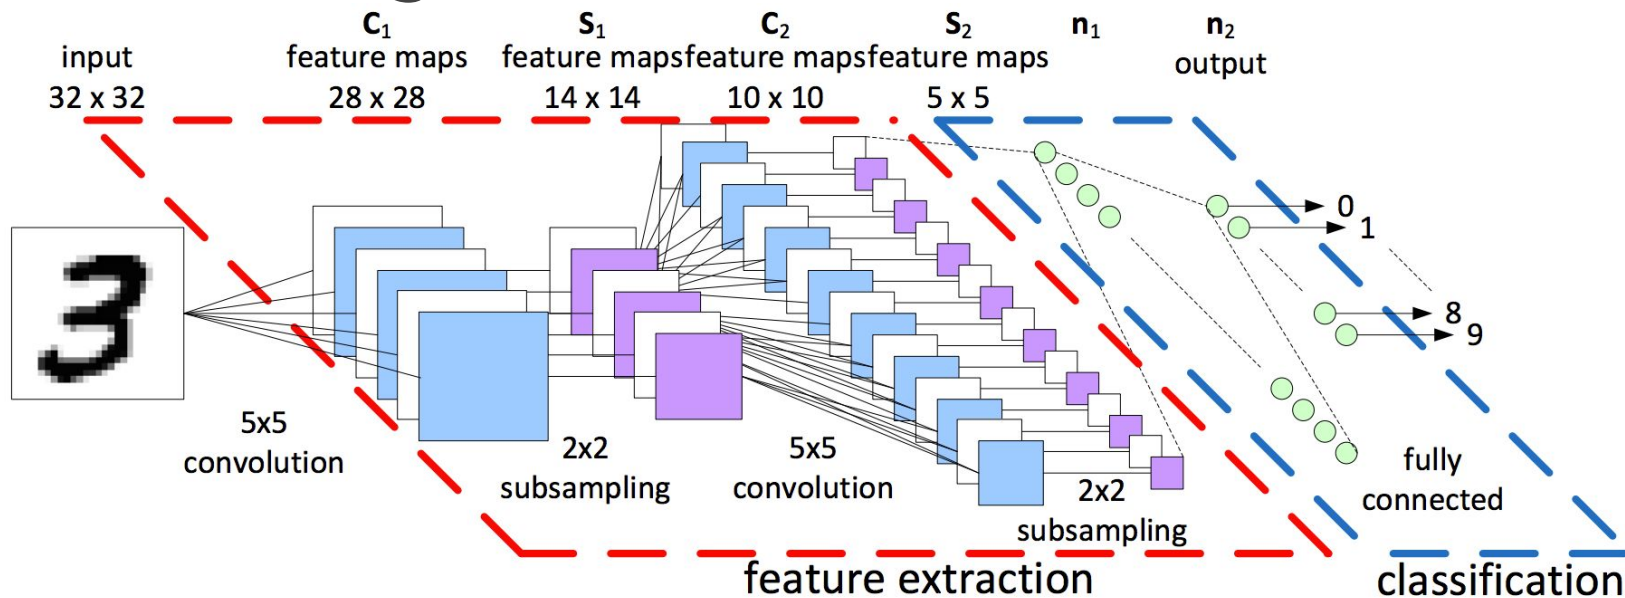
Deep Learning

- This idea then inspired an ANN architecture that would become CNN
- Famously implemented in the 1998 paper by Yann LeCun et al.
- The LeNet-5 architecture was first used to classify the MNIST data set.



Deep Learning

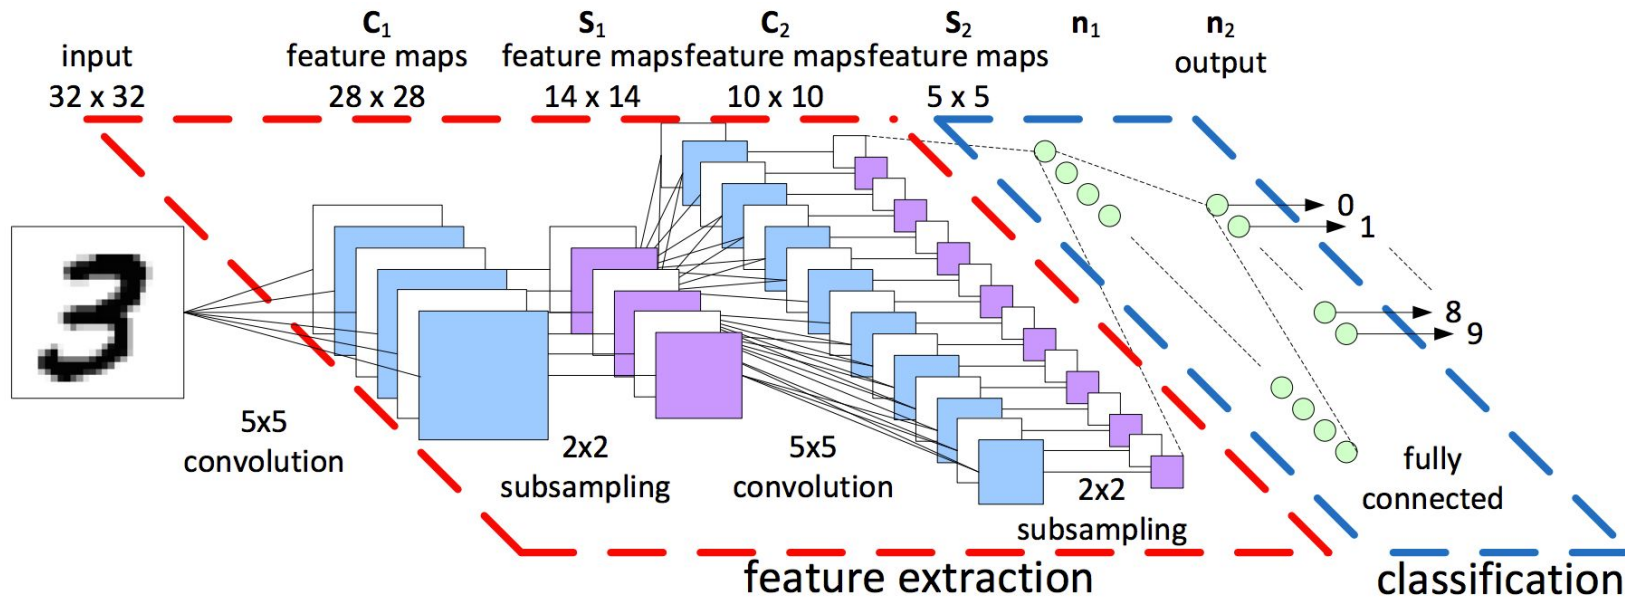
- When learning about CNNs you'll often see a diagram like this:





Deep Learning

- Let's break down the various aspects of a CNN seen here:





Deep Learning

- Tensors
- DNN vs CNN
- Convolutions and Filters
- Padding
- Pooling Layers
- Review Dropout



Deep Learning

- Recall that Tensors are N-Dimensional Arrays that we build up to:
 - Scalar - 3
 - Vector - [3,4,5]
 - Matrix - [[3,4] , [5,6] , [7,8]]
 - Tensor - [[[1, 2] , [3, 4]],
[[5, 6] , [7, 8]]]



Deep Learning

- Tensors make it very convenient to feed in sets of images into our model - (I, H, W, C)
 - I : Images
 - H : Height of Image in Pixels
 - W : Width of Image in Pixels
 - C : Color Channels: 1-Grayscale,



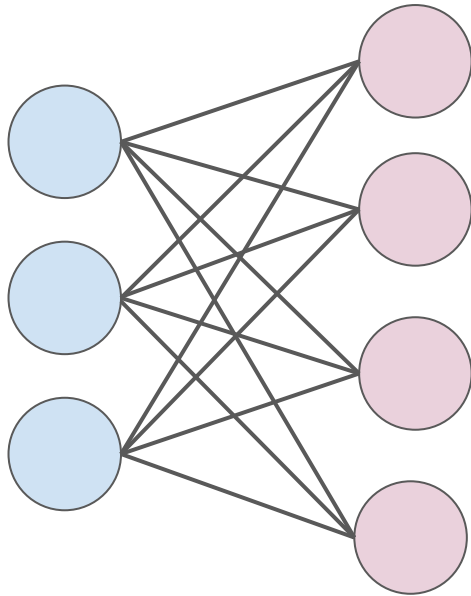
Deep Learning

- Now let's explore the difference between a Densely Connected Neural Network and a Convolutional Neural Network.
- Recall that we've already been able to create DNNs with tf.estimator API.



Deep Learning


- Densely Connected layer





Deep Learning


- Convolutional Layer





Deep Learning

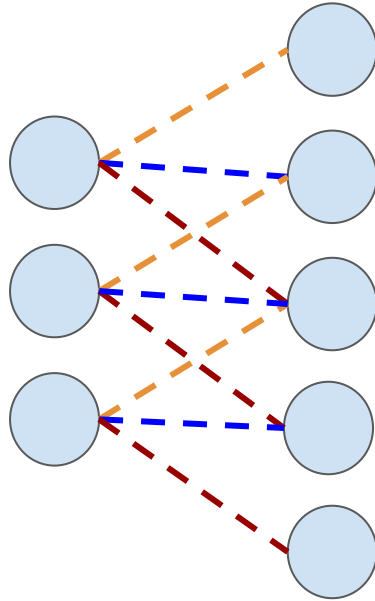
- Each unit is connected to a smaller number of nearby units in next layer.





Deep Learning

- So why bother with a CNN instead of a DNN?






Deep Learning

- The MNIST dataset was 28 by 28 pixels (784 total)
- But most images are at least 256 by 256 or greater, (<56k total)!
- This leads to too many parameters, unscalable to new images.



Deep Learning

- Convolutions also have a major advantage for image processing, where pixels nearby to each other are much more correlated to each other for image detection.





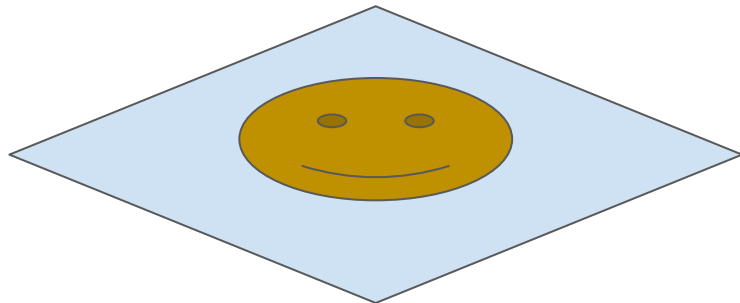
Deep Learning

- Each CNN layer looks at an increasingly larger part of the image.
- Having units only connected to nearby units also aids in *invariance*.
- CNN also helps with regularization, limiting the search of weights to the size of the convolution.



Deep Learning

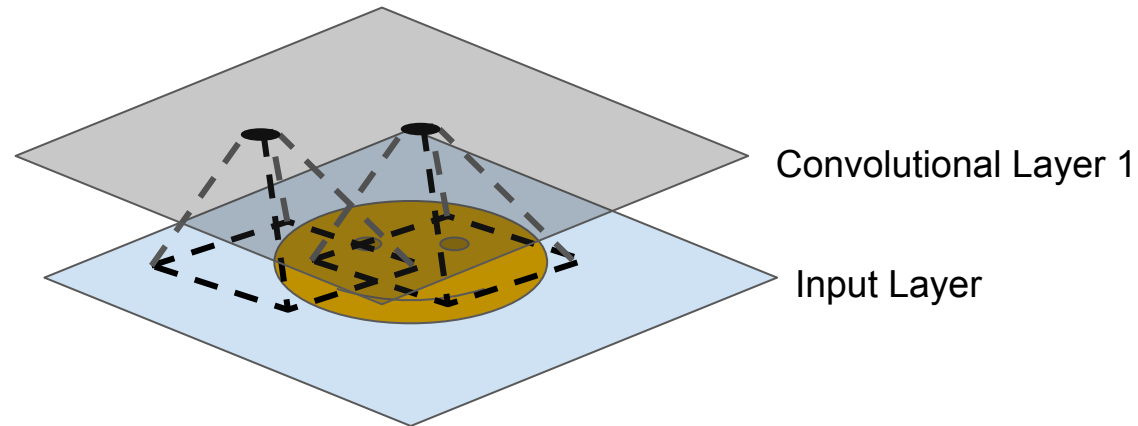
- Let's explore how the convolutional neural network relates to image recognition!
- We start with the input layer, the image itself.





Deep Learning


- Convolutional layers are only connected to pixels in their respective fields.





Deep Learning


- Convolutional layers are only connected to pixels in their respective fields.





Deep Learning

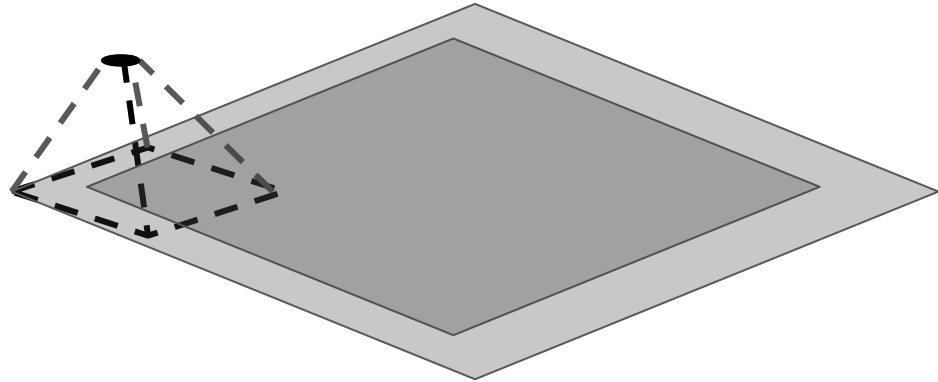
- We run into a possible issue for edge neurons! There may not be an input there for them.





Deep Learning

- We can fix this by adding a “padding” of zeros around the image.





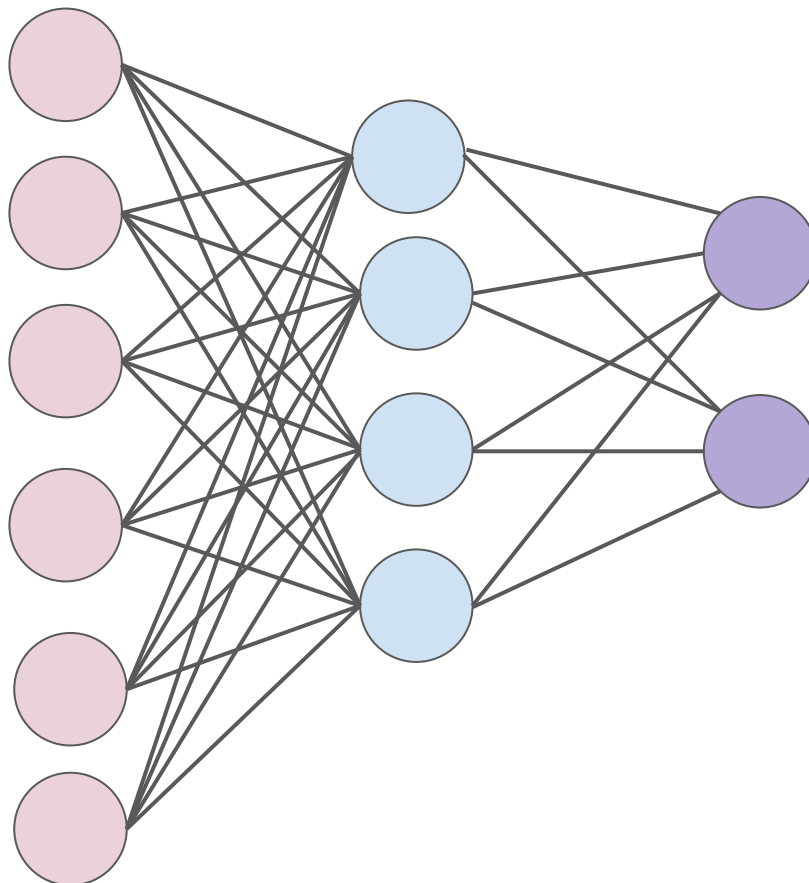
Deep Learning

- Let's walk through 1-D Convolution in more detail, then expand this idea to 2-D Convolution.
- Let's revisit our DNN and convert it to a CNN.



Deep Learning

- A DNN






Deep Learning

1-D

Convolution






Deep Learning

1-D Convolution

We can treat these weights as a filter.





Deep Learning

1-D Convolution


$$y = w_1x_1 + w_2x_2$$

If $(w_1, w_2) = (1, -1)$ Then

$$y = x_1 - x_2$$

When is y at a maximum?

$$(x_1, x_2) = (1, 0)$$





Deep Learning

1-D Convolution


$$y = w_1x_1 + w_2x_2$$

If $(w_1, w_2) = (1, -1)$ Then

$$y = x_1 - x_2$$

When is y at a maximum?

$$(x_1, x_2) = (1, 0)$$






Deep Learning

1-D Convolution

We now have a set of weights that can act as a filter for edge detection!

We can then expand this idea to multiple filters.






Deep Learning

1-D Convolution

Filters: 1

Filter Size: 2

Stride: 2






Deep Learning

1-D Convolution

Filters: 1

Filter Size: 2

Stride: 1 (1 Unit at a time)






Deep Learning

1-D Convolution

Remember that we can add zero padding to include more edge pixels.





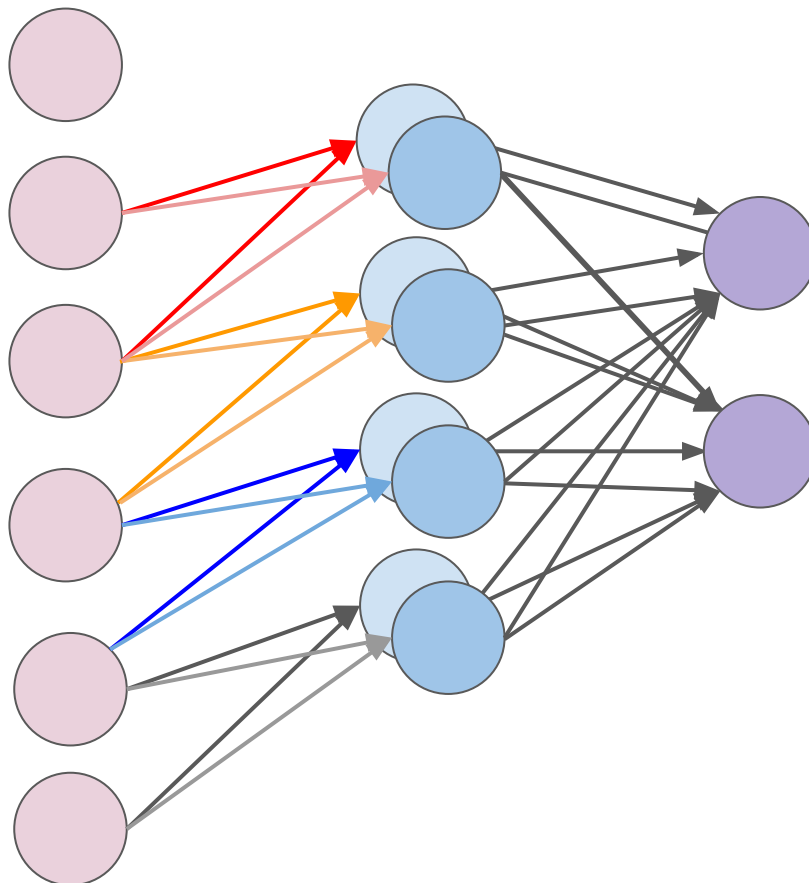
Deep Learning

1-D Convolution

Filters: 2

Filter Size: 2

Stride: 1 (1 Unit at a time)






Deep Learning

1-D Convolution

Filters: 4

Filter Size: 2

Stride: 1 (1 Unit at a time)






Deep Learning

1-D Convolution

Each filter is detecting
a different feature





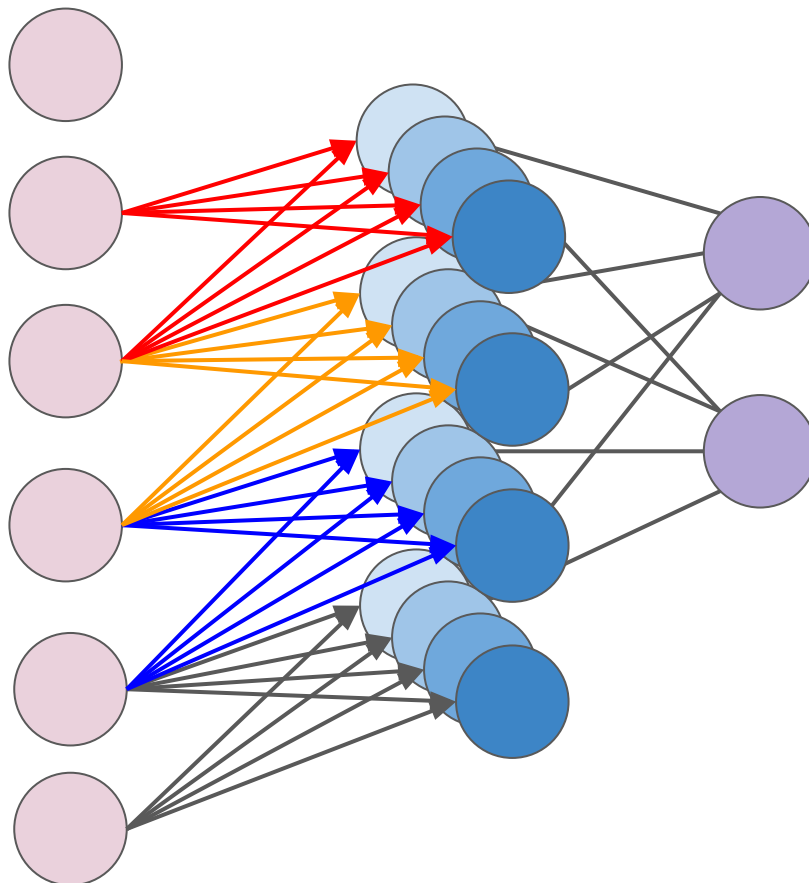
Deep Learning

1-D Convolution

Filters: 4

Filter Size: 2

Stride: 1 (1 Unit at a time)

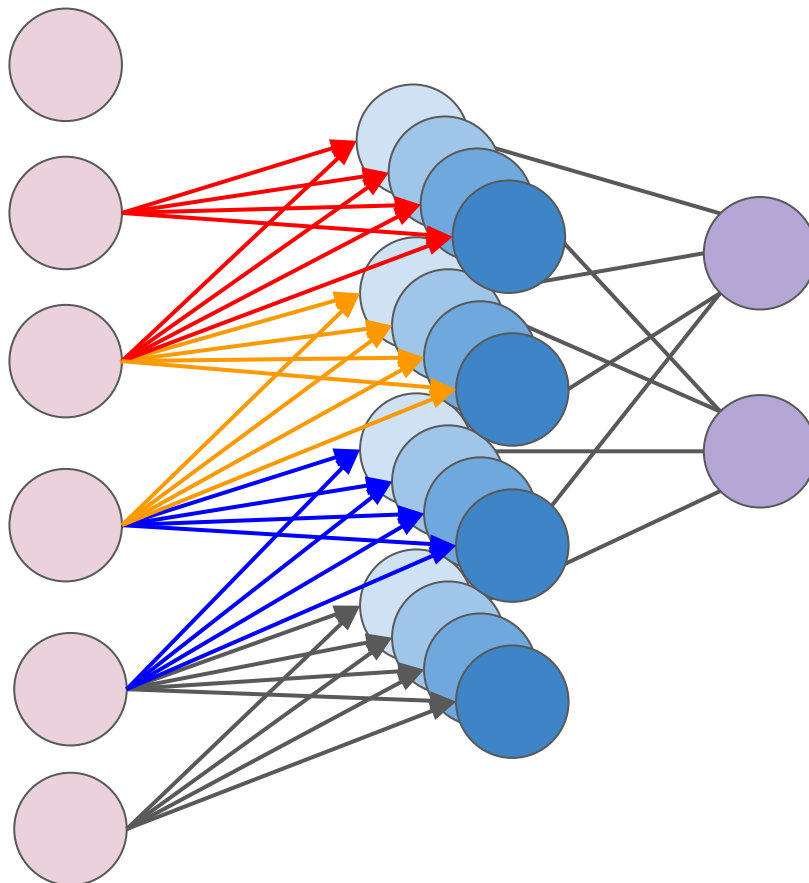




Deep Learning

1-D Convolution

For simplicity, we begin to describe and visualize these sets of neurons as blocks instead







Deep Learning

1-D Convolution

For simplicity, we begin to describe and visualize these sets of neurons as blocks instead



$1 \times L$






Deep Learning

1-D Convolution

For simplicity, we begin to describe and visualize these sets of neurons as blocks instead





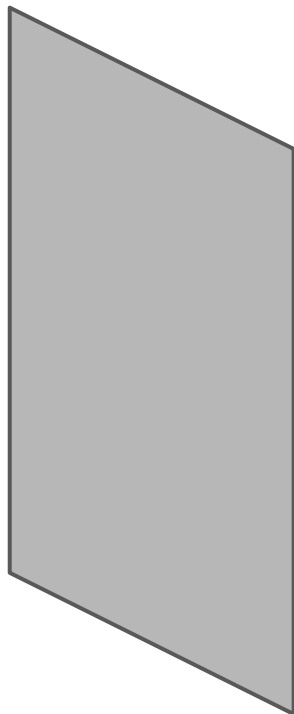
Deep Learning

- Let's now expand these concepts to 2-D Convolution, since we'll mainly be dealing with images.




Deep Learning - 2D Convolution

Input Image:
(H,W)



of filters



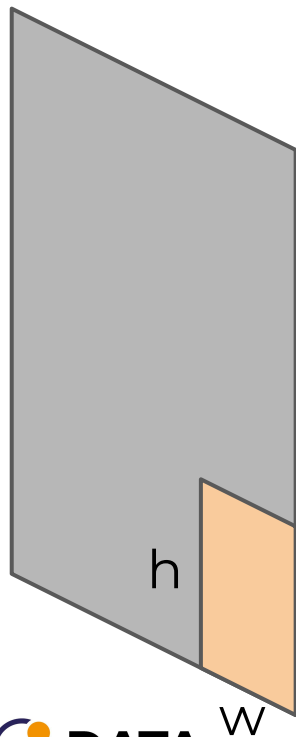
of units W

of units H



Deep Learning - 2D Convolution

Input Image:
(H,W)




of filters



of units W




of units H





Deep Learning - 2D Color Images


Input Image:
(H,W,C)





Deep Learning- 2D Color Images

Input Image:
(H,W,C)





Deep Learning

- Filters are commonly visualized with
grids

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0



Deep Learning

- Filters are commonly visualized with grids

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0


0	0	0
0	-1	1
0	1	-1

3 by 3 Filter



Deep Learning

- Filters are commonly visualized with grids





Deep Learning

- Filters are commonly visualized with grids

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0

0	0	0
0	-1	1
0	1	-1

3 by 3 Filter



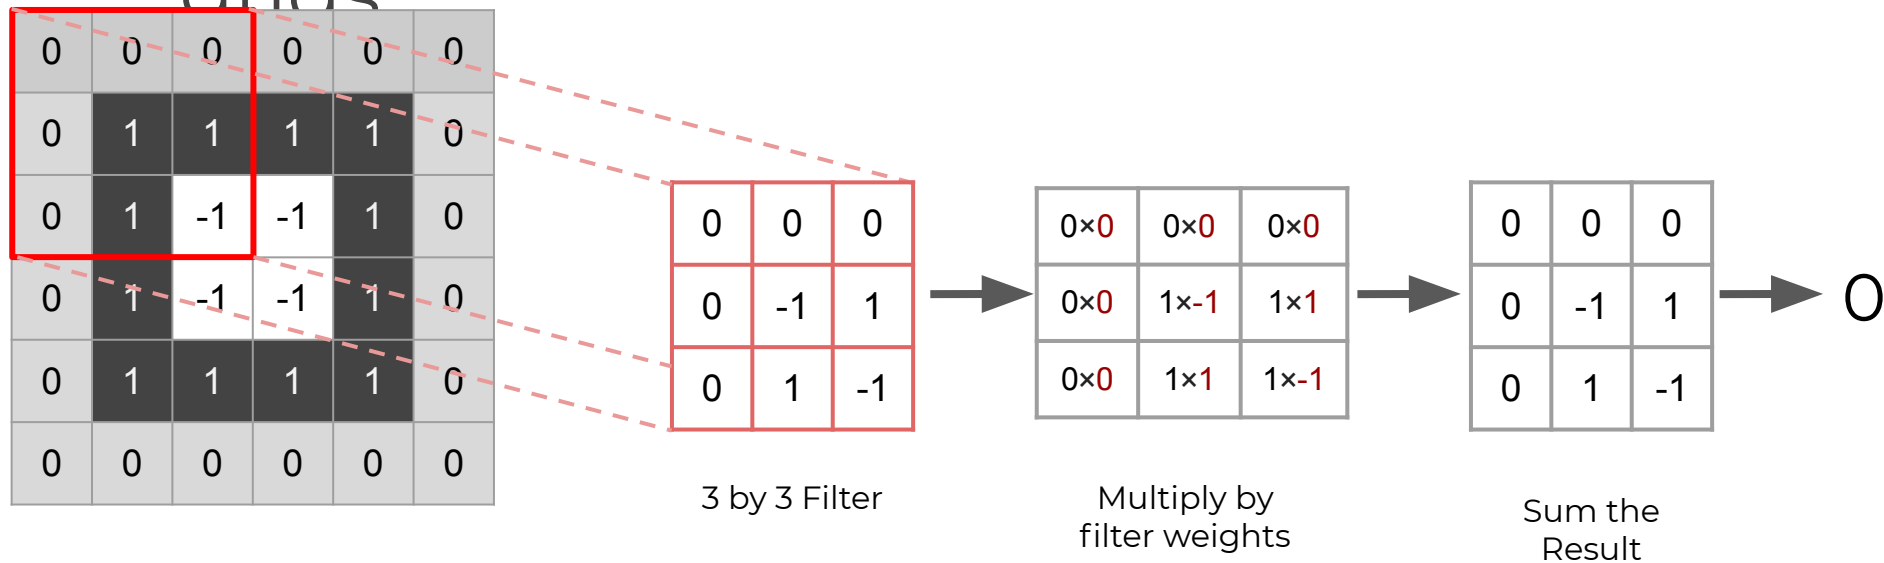
0×0	0×0	0×0
0×0	1×-1	1×1
0×0	1×1	1×-1

Multiply by
filter weights



Deep Learning


- Filters are commonly visualized with grids





Deep Learning


- Stride Distance of 1 Example





Deep Learning

- Stride Distance of 2 Example

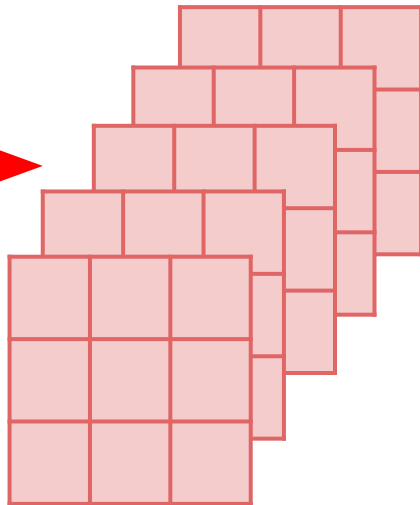




Deep Learning

- Representation of Multiple Filters

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0

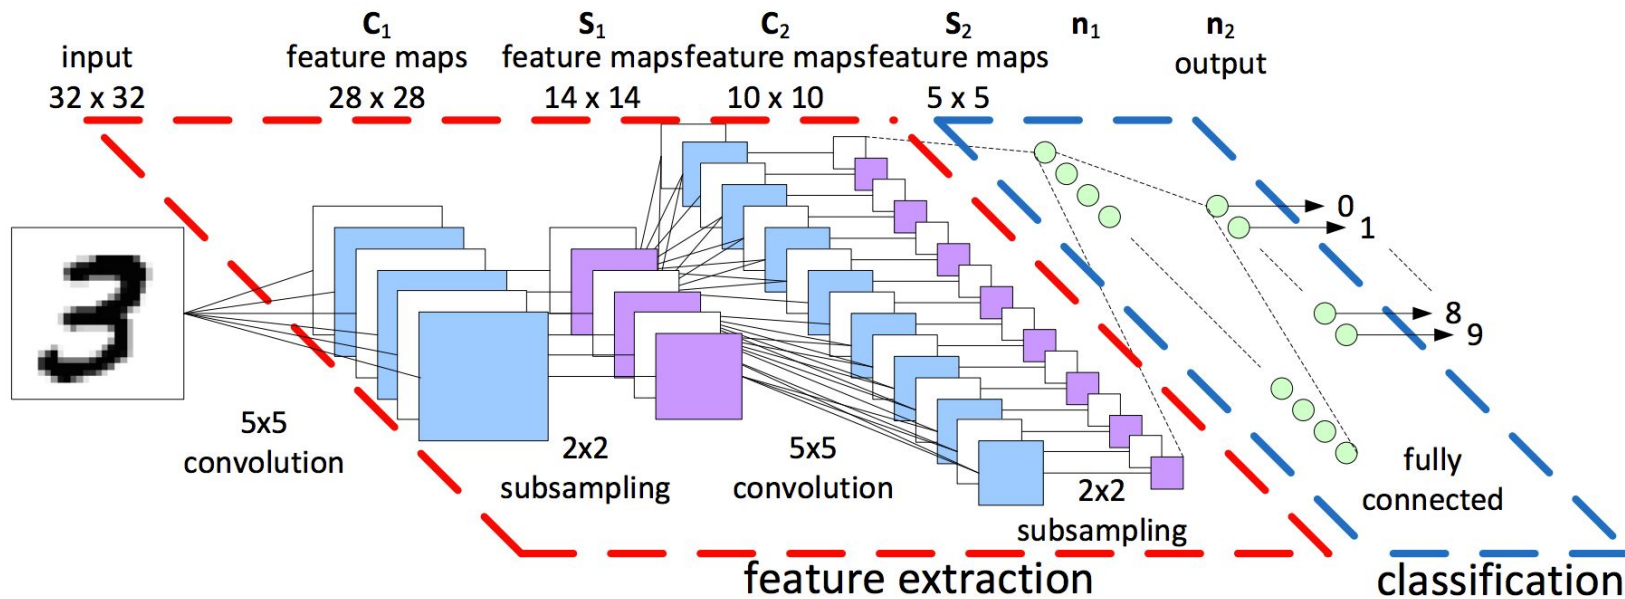


3 by 3
Convolution



Deep Learning


- At the original CNN diagram





Deep Learning


- Exactly what we saw here:





Deep Learning

- Now it is time to discuss subsampling (pooling)





Convolutional Neural Networks Part Two




Deep Learning

- Now that we understand convolutional layers, let's discuss **pooling layers**.
- Pooling layers will subsample the input image, which reduces the memory use and computer load as well as reducing the number of parameters.



Deep Learning

- Let's imagine a layer of pixels in our input image:





Deep Learning

- For our MNIST digits set, each pixel had a value representing “darkness”

0	.2	0	...		
.4	.3	0	...		
...		



Deep Learning


- We create a 2 by 2 pool of pixels and evaluate the maximum value.

0	.2	0	...		
.4	.3	0	...		
...		



Deep Learning

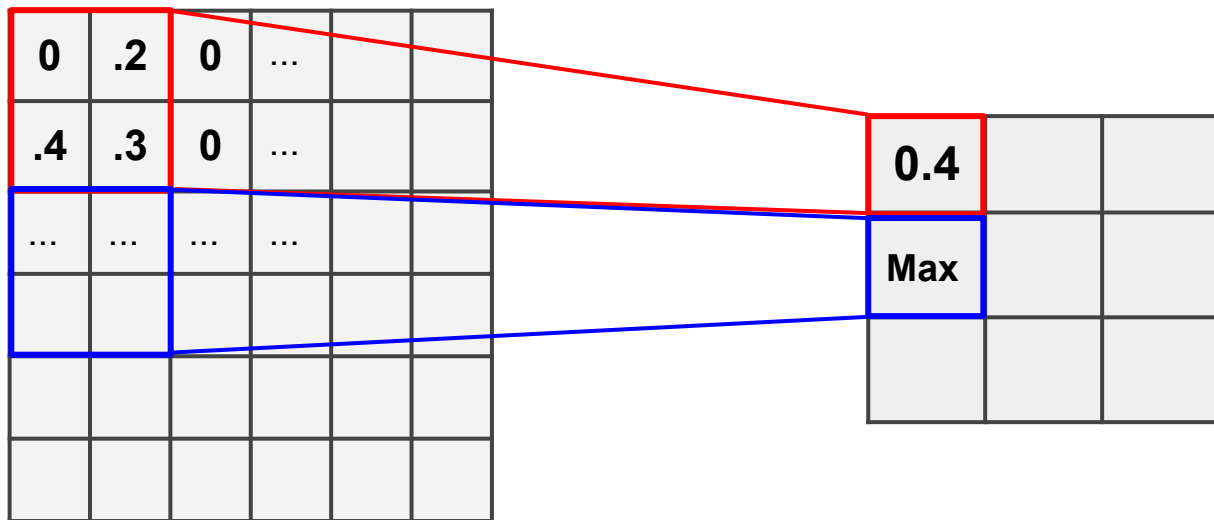
- Only the max value makes it to the next layer.





Deep Learning

- We then move over by a “stride”, in this case, our stride is two.





Deep Learning

- This pooling layer will end up removing a lot of information, even a small pooling “kernel” of 2 by 2 with a stride of 2 will remove 75% of the input data.



Deep Learning

- Another common technique deployed with CNN is called “Dropout”
- Dropout can be thought of as a form of regularization to help prevent overfitting.
- During training, units are randomly dropped, along with their connections.



Deep Learning

- This helps prevent units from “co-adapting” too much.
- Let’s also quickly point out some famous CNN architectures!



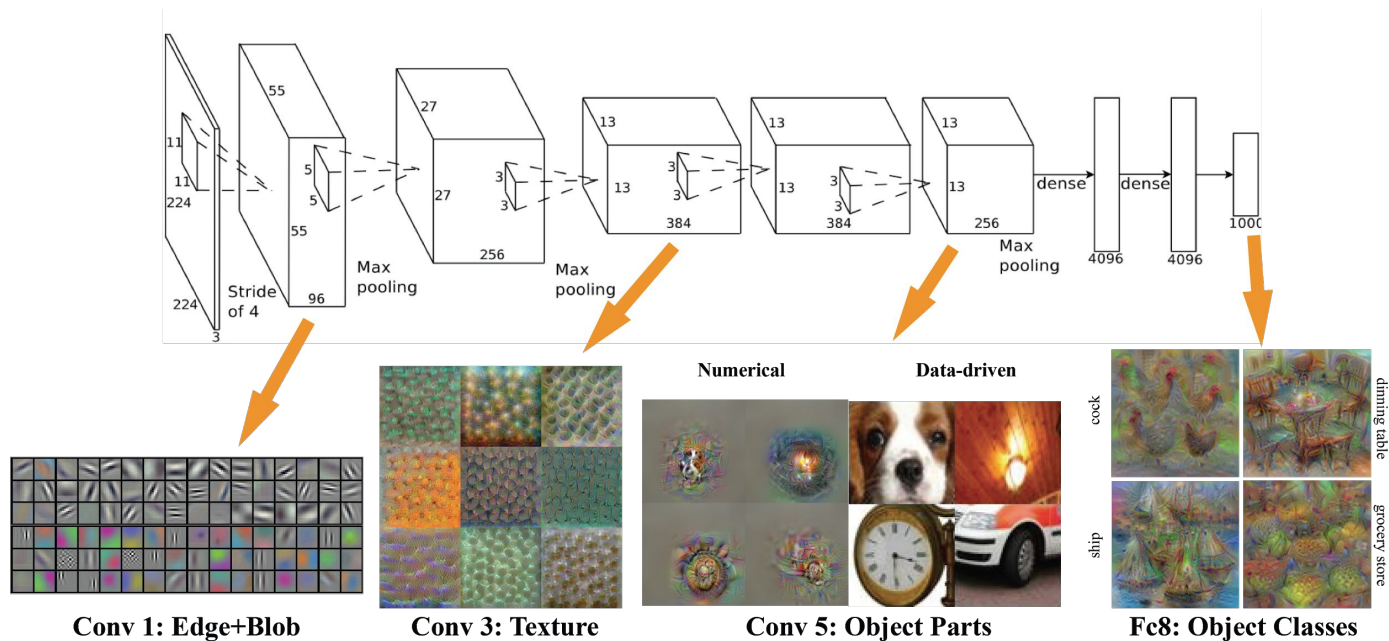
Deep Learning

- LeNet-5 by Yann LeCun
- AlexNet by Alex Krizhevsky et al.
- GoogLeNet by Szegedy at Google Research
- ResNet by Kaiming He et al.
- Check out the resource links to the papers discussing these architectures!



Deep Learning


- AlexNet





Deep Learning

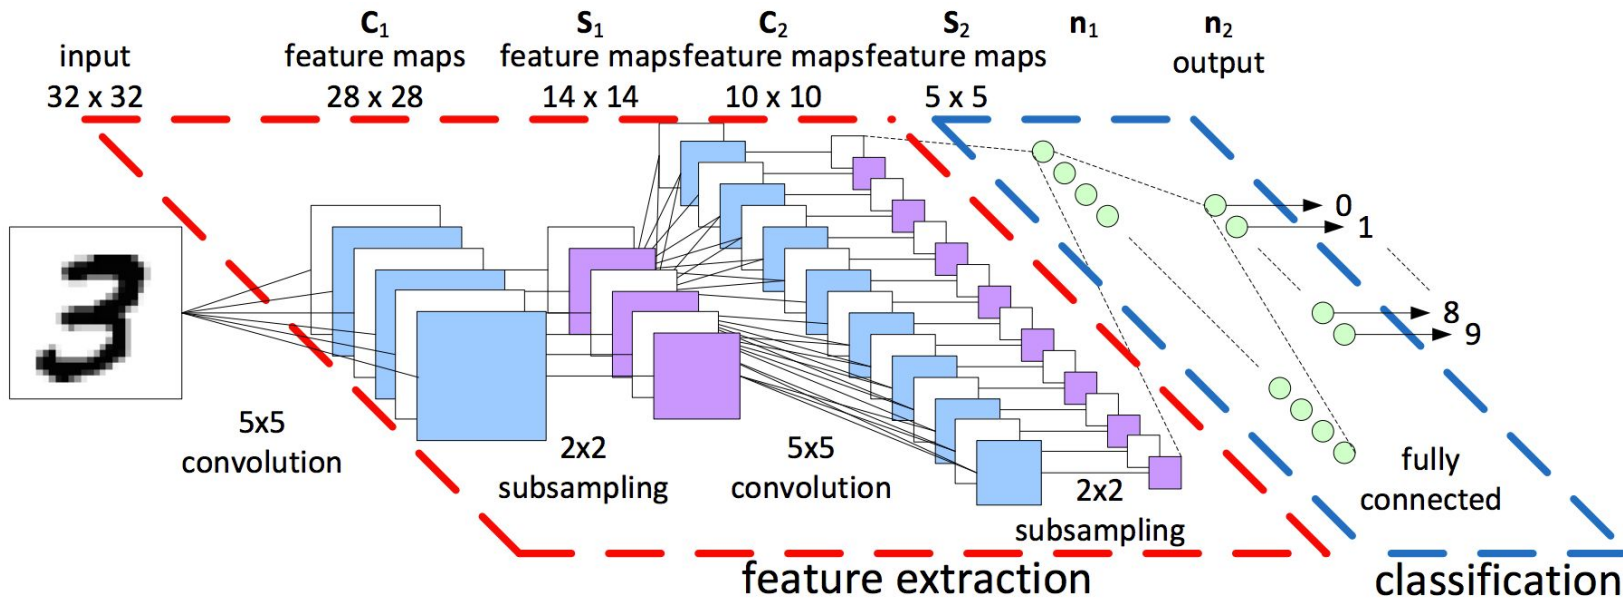
- AlexNet





Deep Learning

- Convolutional Neural Network





Deep Learning

- Check out the various supplementary resources!
- Now that we've covered the basics of CNN, let's explore how to implement one in TensorFlow!



MNIST Data CNN Code Along



CNN CIFAR-10 Exercise



CNN CIFAR-10 Exercise Solutions Part Two




Deep Learning

- Optional CNN Exercise with the CIFAR-10 data set.
- Main challenge is dealing with data and creating Tensor batches and sizing.
- Let's go through the exercise notebook!



Deep Learning

- We have our color image represented as a 3-D prism. Width, Height, and Depth (3 values R, G, B)





Deep Learning

- Run a neural network on a portion

